

SOLUTIONS

EXERCISE 1

Parameter combination table using the abalone data set:

```
> accuracies_table
  Degree Cost CV.Accuracy Entire.DF.Accuracy
4      1 100.0    26.21499         27.65142
3      1  10.0    25.78406         26.88532
2      1   1.0    25.30524         26.19105
1      1  0.1    24.01245         24.63491
8      2 100.0    26.64592         31.79315
7      2  10.0    26.50227         28.82452
6      2   1.0    25.97558         27.43596
5      2  0.1    23.94063         24.87431
12     3 100.0    26.64592         33.68446
11     3  10.0    26.09528         29.56667
10     3   1.0    24.97007         26.71774
9      3  0.1    23.53364         24.61096
```

Combination with the highest cross-validation accuracy:

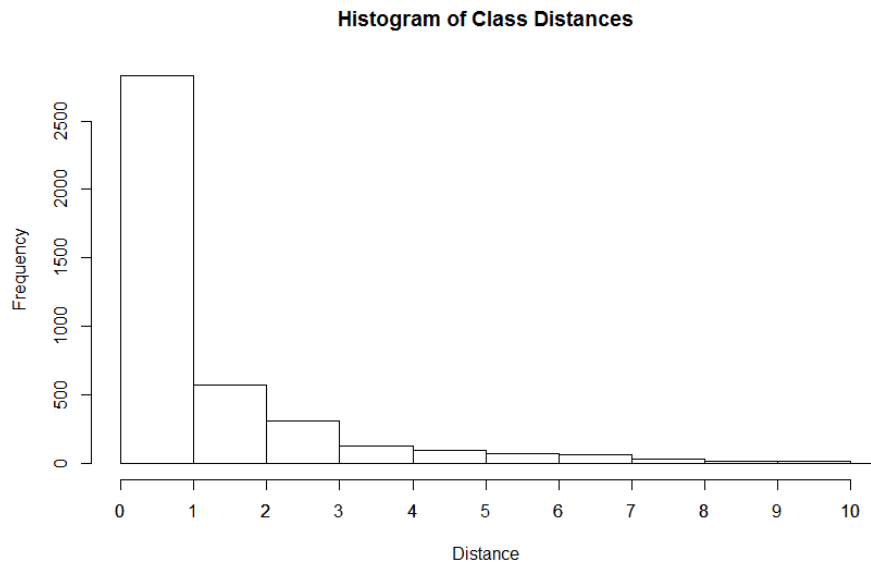
```
> max_combination
  Degree Cost CV.Accuracy Entire.DF.Accuracy
8      2  100    26.64592         31.79315
```

We see that the degree and cost values that produce the high cross-validation accuracy, 26.6, are 2 and 100, respectively.

Average distance of the predicted class from the true class:

```
> average_distance
[1] 1.513766
```

Histogram showing the frequency of how often a prediction is m away from the true number of rings:



EXERCISE 2

Table of best parameter combinations for each trained binary classifier (trained using abalone data set):

```
> binary_classifier_table
```

| | Description | Dataset.Size | Degree | Cost | Average.CV.Accuracy | Best.Accuracy |
|---|------------------|--------------|--------|------|---------------------|------------------|
| 1 | <= 9 vs >= 10 | 4177 | 2 | 100 | 77.0868246748065 | 79.8659324874312 |
| 2 | <= 7 vs 8 - 9 | 2096 | 1 | 10 | 82.0928753180662 | 83.0629770992366 |
| 3 | <= 5 vs 6 - 7 | 839 | 1 | 0.1 | 87.0977353992849 | 87.7234803337306 |
| 4 | 8 vs 9 | 1257 | 2 | 1 | 62.9209758684699 | 64.7573587907717 |
| 5 | 6 vs 7 | 650 | 3 | 1 | 64.3076923076923 | 65.8461538461538 |
| 6 | 10 - 11 vs >= 12 | 2081 | 2 | 100 | 65.8537562069518 | 70.783277270543 |
| 7 | 12 - 13 vs >= 14 | 960 | 1 | 10 | 60.1388888888889 | 64.375 |
| 8 | 10 vs 11 | 1121 | 3 | 10 | 57.8873030032709 | 59.2328278322926 |
| 9 | 12 vs 13 | 470 | 2 | 100 | 56.8439716312057 | 59.5744680851064 |

EXERCISE 3

Average distance of the predicted class from the true class of the binary-search learning algorithm applied to the abalone data set:


```
> binary_average_distance
[1] 1.141
```

It is worth noting that the average distance using the binary search learning algorithm is about 0.4 lower.

Training accuracy of the binary-search learning algorithm applied to the abalone data set:

```
> binary_class_accuracy
[1] 36.62916
```

Histogram showing the frequency of how often a prediction is m rings away from the true value:



EXERCISE 4

Table of 35 parameter combination results for eps-regression performed on the Exercise-4 data set:

```
> reg_accuracies_table
  Epsilon Cost    CV.MSE Entire.DF.MSE
1    0.25 1e-01  1.957467    1.954661
2    0.50 1e-01  4.841288    4.635268
3    0.75 1e-01 10.620081   10.494145
4    1.00 1e-01 20.191924   19.753267
5    1.25 1e-01 31.890114   31.327224
6    1.50 1e-01 45.605031   45.503308
7    1.75 1e-01 64.058366   63.835242
8    0.25 1e+00  1.901948    1.880863
9    0.50 1e+00  3.009844    2.937969
10   0.75 1e+00  9.455004    9.389761
11   1.00 1e+00 18.188184   18.107956
12   1.25 1e+00 30.246685   30.152740
13   1.50 1e+00 45.629492   45.485435
14   1.75 1e+00 64.005422   63.828593
15   0.25 1e+01  1.857911    1.860247
16   0.50 1e+01  3.087263    2.941775
17   0.75 1e+01  9.435075    9.388801
18   1.00 1e+01 18.153996   18.114872
19   1.25 1e+01 30.258712   30.153416
20   1.50 1e+01 45.669981   45.492341
21   1.75 1e+01 63.988900   63.843066
22   0.25 1e+02  1.870768    1.868057
23   0.50 1e+02  3.070670    2.942062
24   0.75 1e+02  9.412851    9.390340
25   1.00 1e+02 18.184213   18.116082
26   1.25 1e+02 30.268778   30.152287
27   1.50 1e+02 45.597345   45.486227
28   1.75 1e+02 63.954225   63.826524
29   0.25 1e+03  1.861971    1.853047
30   0.50 1e+03  3.011352    2.946526
31   0.75 1e+03  9.424377    9.389048
32   1.00 1e+03 18.204606   18.116458
33   1.25 1e+03 30.275061   30.153726
34   1.50 1e+03 45.609684   45.506906
35   1.75 1e+03 64.016537   63.824236
```

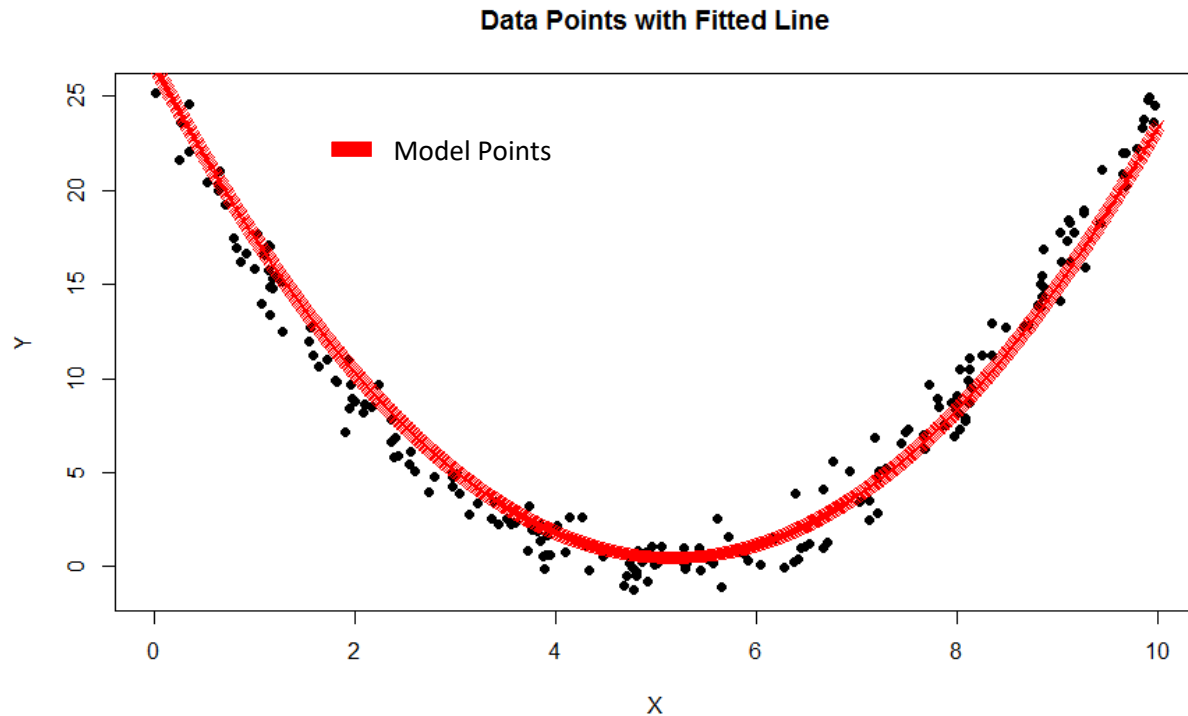
Combination resulting in the lowest* average CV MSE:

```
> reg_min_combination
  Epsilon Cost    CV.MSE Entire.DF.MSE
15    0.25   10  1.857911    1.860247
```

We see that the epsilon and cost values that produce the lowest error, 1.857911, are 0.25 and 10, respectively.

EXERCISE 5

Graph of plotted data points against curve provided by the best SVM from the previous exercise:



The **predicted line** was found by creating a sequence of 1000 points evenly spaced between 0 and 10.

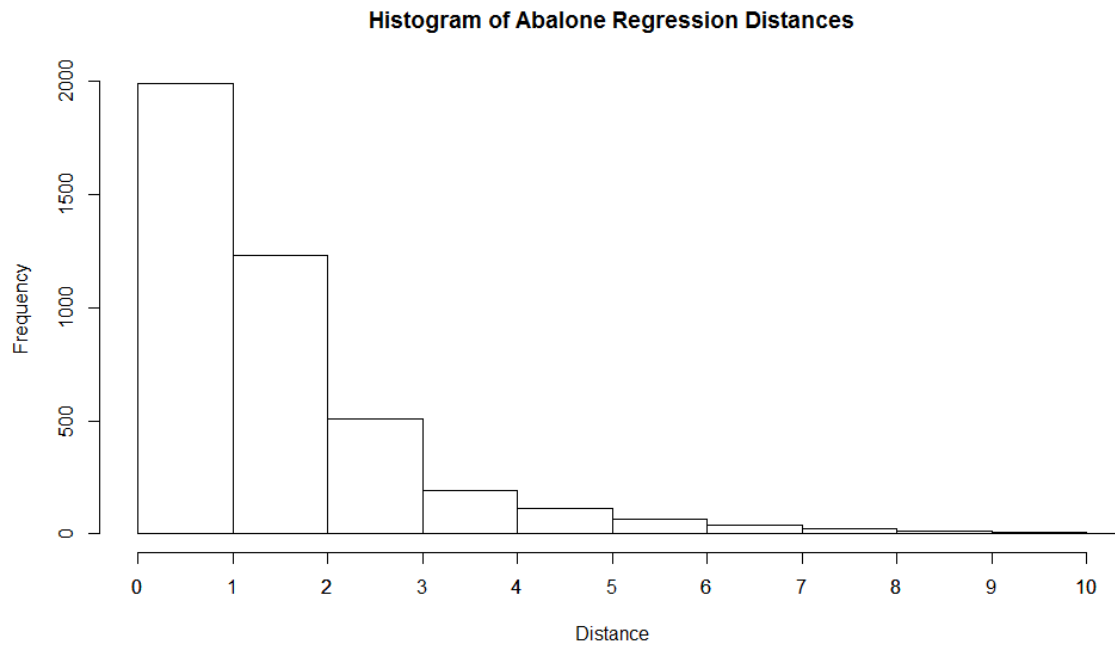
EXERCISE 6

Average distance of the predicted class from the true class using the d , C , and e parameters which produce the least MSE:

```
> average_distance_reg_abalone  
[1] 1.462579
```

Just to be clear, a table was not asked for in this exercise.

Histogram showing frequency of how often a prediction is m rings away from the true number of rings:



APPENDIX

EXERCISE 1

EXERCISE 1

READ IN ABALONE DATA SET AND SET THE COLUMN NAMES BASED ON DATA DESCRIPTION

```
abalone_df <- read.table("abalone.data", sep = ",", header = FALSE)
colnames(abalone_df) <- c("Sex", "Length", "Diam", "Height", "Whole",
  "Shucked", "Viscera", "Shell", "Rings")
```

INITIALIZE LEARNING PARAMETERS

```
degrees <- c(1:3)
costs <- c(10 ^ (-1:2))
cross_fold <- 5
```

KEEP TRACK OF BEST CV ACCURACY AND THE BEST RESPECTIVE PREDICTIONS

```
best_cv_accuracy <- 0
best_predictions <- data.frame()
```

INITIALIZE TABLE (DATA FRAME TO BE APPENDED)

```
accuracies_table <- data.frame("Degree" = integer(), "Cost" = numeric(),
  "CV Accuracy" = numeric(), "Entire DF Accuracy" = numeric(),
  stringsAsFactors = FALSE)
```

BUILD MODEL FOR EVERY COMBINATION OF

```
for (d in degrees) {
  for (c in costs) {
```

MODEL WITHOUT CROSS VALIDATION

```
current_model_entire_df <- svm(Rings ~ ., data = abalone_df,
  kernel = "polynomial", degree = d, type = "C-classification", cost = c)
```

```
predictions <- predict(current_model_entire_df, abalone_df[, - length(abalone_df)])
accuracy_total <- 100 * mean(predictions == abalone_df[, length(abalone_df)])
```

MODEL WITH 5-FOLD CROSS VALIDATION

```
current_model <- svm(Rings ~ ., data = abalone_df, kernel = "polynomial",
  degree = d, type = "C-classification", cost = c, cross = cross_fold)
```

```
accuracies_table[nrow(accuracies_table) + 1,] <- c(d, c, current_model$tot.accuracy,
accuracy_total)
```

```
if (current_model$tot.accuracy > best_cv_accuracy) {
  best_cv_accuracy <- current_model$tot.accuracy
  best_predictions <- predict(current_model, abalone_df[, -length(abalone_df)])
}
```

```
}
}
```

SORT BY INCREASING COMPLEXITY

```
accuracies_table <- accuracies_table[order(-accuracies_table$Cost),]
accuracies_table <- accuracies_table[order(accuracies_table$Degree),]
```

FIND COMBINATION WITH HIGHEST CV ACCURACY

```
max_combination <- accuracies_table[which.max(accuracies_table$CV.Accuracy),]
```

AVERAGE DISTANCE OF THE PREDICTED CLASS FROM THE TRUE CLASS USING BEST PARAMS

```

best_pred_int <- as.integer(best_predictions)
class_distances <- abs(best_pred_int - abalone_df$Rings)
average_distance <- mean(class_distances)

# HISTOGRAM SHOWING FREQUENCY OF M DISTANCES FROM TRUE CLASS
par(bg = 'white')
hist(class_distances, main = "Histogram of Class Distances", xlab = "Distance", ylab = "Frequency",
xlim = c(0, 10))
axis(side=1, at=seq(0,10,1), labels=seq(0,10,1))

```

EXERCISE 2

#EXERCISE 2

```

# MAKE ALL LESS THAN 5 = 5
abalone_df$Rings[abalone_df$Rings < 5] <- 5
# MAKE ALL GREATER THAN 14 = 14
abalone_df$Rings[abalone_df$Rings > 14] <- 14

# INITIALIZE BOUNDS FOR BINARY CLASSIFIER
f1 <- list(c(0:9), c(10:30))
f2 <- list(c(0:7), c(8:9))
f3 <- list(c(0:5), c(6:7))
f4 <- list(c(8), c(9))
f5 <- list(c(6), c(7))
f6 <- list(c(10:11), c(12:30))
f7 <- list(c(12:13), c(14:30))
f8 <- list(c(10), c(11))
f9 <- list(c(12), c(13))

# LIST OF ALL BOUNDS
classifier_bounds <- list(f1, f2, f3, f4, f5, f6, f7, f8, f9)

# TABLE (DATA FRAME) WHERE CLASSIFIER TRAINING RESULTS WILL RESIDE
binary_classifier_table <- data.frame("Description" = character(), "Dataset Size" = integer(), "Degree"
= numeric(), "Cost" = numeric(),
                                     "Average CV Accuracy" = numeric(), "Best Accuracy" = numeric(),
stringsAsFactors = FALSE)

# LIST OF ALL QUERY NODE MODELS
binary_classifier_models <- list()

for (bound in classifier_bounds) {

  # CALCULATE BEST PARAMATERS GIVEN THE BINARY PARTITION BOUNDS AND APPEND TO CLASSIFIER PARAMS
  DATAFRAME
  negative_class <- bound[[1]]
  positive_class <- bound[[2]]

  # CREATE DESCRIPTION BASED ON BOUNDS
  description <- ""
  if (length(negative_class) > 2) {
    description <- paste(description, "<=", negative_class[length(negative_class)])
  } else if (length(negative_class) == 2) {
    description <- paste(description, negative_class[1], "-", negative_class[2])
  } else {
    description <- paste(description, negative_class[1])
  }
  description <- paste(description, " vs ")
  if (length(positive_class) > 2) {
    description <- paste(description, ">=", positive_class[1])
  }
}

```

```

} else if (length(positive_class) == 2) {
  description <- paste(description, positive_class[1], "-", positive_class[2])
} else {
  description <- paste(description, positive_class[1])
}

df <- abalone_df

# CREATE SUBSET BASED ON BOUNDS
df <- df[df$Rings %in% negative_class | df$Rings %in% positive_class,]

# CHANGE PROBLEM TO A BINARY CLASSIFIER BASED ON BOUNDS
df$Rings[df$Rings %in% negative_class] <- -1
df$Rings[df$Rings %in% positive_class] <- 1

average_accuracy <- 0
best_accuracy <- 0
best_d <- 0
best_c <- 0
cross_fold <- 5
num_of_combos <- length(costs) * length(degrees)
best_model <- 0

for (d in degrees) {
  for (c in costs) {

    # BUILD MODEL FOR GIVEN COMBINATION OF D AND C
    current_model <- svm(Rings ~ ., data = df, kernel = "polynomial",
      degree = d, type = "C-classification", cost = c, cross = cross_fold)
    average_accuracy = average_accuracy + current_model$tot.accuracy

    # FIND BEST ACCURACY
    if (current_model$tot.accuracy > best_accuracy) {
      best_accuracy <- current_model$tot.accuracy
      best_d <- d
      best_c <- c
      best_model <- current_model
    }
  }
}

# APPEND BINARY CLASSIFIER TABLE: DESCRIPTION, TRAINING SET SIZE, BEST D, BEST C, AVERAGE CV ACC,
BEST ACC
binary_classifier_table[nrow(binary_classifier_table) + 1,] <- c(description, nrow(df), best_d,
best_c, average_accuracy / num_of_combos, best_accuracy)

# APPEND BINARY CLASSIFIER MODEL LIST: THESE ARE THE QUERY NODES
binary_classifier_models[[length(binary_classifier_models) + 1]] <- best_model
}

```

EXERCISE 3

EXERCISE 3

PREDICT CLASSIFICATIONS USING TRAINED BINARY CLASSIFIER QUERA NODES

```

final_predictions <- c()

for (i in 1:nrow(abalone_df)) {
  # <= 9 or >= 10

```



```
model <- binary_classifier_models[[1]]
prediction <- predict(model, abalone_df[i, - length(abalone_df)])

if (prediction == 1) {
  # prediction >= 10

  # 10-11 or >= 12
  model <- binary_classifier_models[[6]]
  prediction <- predict(model, abalone_df[i, - length(abalone_df)])

  if (prediction == 1) {
    # prediction >= 12

    # 12-13 or >= 14
    model <- binary_classifier_models[[7]]
    prediction <- predict(model, abalone_df[i, - length(abalone_df)])

    if (prediction == 1) {
      # final prediction is 14
      final_predictions <- c(final_predictions, 14)
    } else {
      # 12 or 13
      model <- binary_classifier_models[[9]]
      prediction <- predict(model, abalone_df[i, - length(abalone_df)])

      if (prediction == 1) {
        # final prediction is 13
        final_predictions <- c(final_predictions, 13)
      } else {
        # final prediction is 12
        final_predictions <- c(final_predictions, 12)
      }
    }
  }
} else {
  # prediction 10-11

  model <- binary_classifier_models[[8]]
  prediction <- predict(model, abalone_df[i, - length(abalone_df)])

  if (prediction == 1) {
    # final prediction is 11
    final_predictions <- c(final_predictions, 11)
  } else {
    # final prediction is 10
    final_predictions <- c(final_predictions, 10)
  }
}

} else if (prediction == -1){
  # prediction <= 9

  # <= 7 or 8-9
  model <- binary_classifier_models[[2]]
  prediction <- predict(model, abalone_df[i, - length(abalone_df)])

  if (prediction == 1) {
    # prediction 8-9

    model <- binary_classifier_models[[4]]
```

```

prediction <- predict(model, abalone_df[i, - length(abalone_df)])

if (prediction == 1) {
  # final prediction is 9
  final_predictions <- c(final_predictions, 9)
} else {
  # final prediction is 8
  final_predictions <- c(final_predictions, 8)
}

} else {
  # prediction <= 7

  # <= 5 or 6-7
  model <- binary_classifier_models[[3]]
  prediction <- predict(model, abalone_df[i, - length(abalone_df)])

  if (prediction == 1) {
    # prediction 6-7

    model <- binary_classifier_models[[5]]
    prediction <- predict(model, abalone_df[i, - length(abalone_df)])

    if (prediction == 1) {
      # final prediction is 7
      final_predictions <- c(final_predictions, 7)
    } else {
      # final prediction is 6
      final_predictions <- c(final_predictions, 6)
    }
  } else {
    # final prediction is 5
    final_predictions <- c(final_predictions, 5)
  }
}

}

}

# TRAINING ACCURACY OF BINARY CLASSIFIER
binary_class_accuracy <- mean(final_predictions == abalone_df$Rings) * 100

# AVERAGE DISTANCE OF PREDICTED CLASS FROM TRUE CLASS
binary_class_distances <- abs(final_predictions - abalone_df$Rings)
binary_average_distance <- mean(binary_class_distances)

hist(binary_class_distances, main = "Histogram of Binary Class Distances",
      xlab = "Distance", ylab = "Frequency", xlim = c(0, 8))
axis(side = 1, at = seq(0, 10, 1), labels = seq(0, 10, 1))

```

EXERCISE 4

EXERCISE 4

READ IN DATA SET

```
reg_df <- read.csv("Exercise-4.csv")
```

INITIALIZE TRAINING PARAMETERS

```

reg_cross_fold <- 10
reg_costs <- c(10 ^ (-1:3))
epsilons <- c(0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75)

# KEEP TRACK OF PARAMS THAT PRODUCE LOWEST MSE
lowest_mse <- 1000
best_reg_c <- 0
best_reg_e <- 0
best_reg_model <- 0

# TABLE(DATAFRAME) WHERE RESULTS FROM EACH COMBINATION RESIDE
reg_accuracies_table <- data.frame("Epsilon" = integer(), "Cost" = numeric(),
  "CV MSE" = numeric(), "Entire DF MSE" = numeric(), stringsAsFactors = FALSE)

for (c in reg_costs) {
  for (e in epsilons) {

    # BUILD MODEL AND CALCULATE MSE OVER ENTIRE DATASET
    reg_model_entire <- svm(Y ~ X, data = reg_df, kernel = "polynomial",
      degree = 2, type = "eps-regression", epsilon = e, cost = c, cross = 100)

    # BUILD MODEL USING 10-FOLD CROSSVALIDATION
    reg_model <- svm(Y ~ X, data = reg_df, kernel = "polynomial",
      degree = 2, type = "eps-regression", epsilon = e, cost = c, cross = 10)

    # APPEND TABLE WITH CURRENT RESULTS
    reg_accuracies_table[nrow(reg_accuracies_table) + 1,] <- c(e, c, reg_model$tot.MSE,
      reg_model_entire$tot.MSE)

    if (reg_model$tot.MSE < lowest_mse) {
      lowest_mse = reg_model$tot.MSE
      best_reg_c <- c
      best_reg_e <- e
      best_reg_model <- reg_model
    }
  }
}

# SORT TABLE WITH INCREASING COMPLEXITY
reg_accuracies_table <- reg_accuracies_table[order(reg_accuracies_table$Epsilon),]
reg_accuracies_table <- reg_accuracies_table[order(reg_accuracies_table$Cost),]

# COMBINATION FOR HIGHEST CV ACCURACY
reg_min_combination <- reg_accuracies_table[which.min(reg_accuracies_table$CV.MSE),]

```

EXERCISE 5

EXERCISE 5

```

# PREDICTED LINE USING 1000 EVENLY SPACED POINTS FROM 0 TO 10
plot(reg_df, pch = 16)
test_data <- data.frame(seq(0, 10, length.out = 1000))
colnames(test_data) <- c("X")
predicted <- predict(best_reg_model, test_data)
points(test_data$X, predicted, col = "red", pch = 4)
title("Data Points with Fitted Line")

```

EXERCISE 6

EXERCISE 6

INITIALIZE LEARNING PARAMETERS

```
reg_cross_fold <- 5
reg_costs <- c(10 ^ (-1:3))
epsilons <- c(0.25, 0.5, 0.75)
degrees <- c(1:3)
```

KEEP TRACK OF PARAMS THAT PRODUCE LOWEST MSE

```
best_reg_c_abalone <- 0
best_reg_e_abalone <- 0
best_reg_d_abalone <- 0
best_reg_model_abalone <- 0
lowest_mse_abalone <- 1000
best_reg_model_abalone <- 0
```

```
for (c in reg_costs) {
  for (e in epsilons) {
    for (d in degrees) {

      reg_model <- svm(Rings ~ ., abalone_df, kernel = "polynomial", degree = d, type = "eps-
regression", epsilon = e, cost = c, cross = reg_cross_fold)

      if (reg_model$tot.MSE < lowest_mse_abalone) {
        lowest_mse_abalone = reg_model$tot.MSE
        best_reg_c_abalone <- c
        best_reg_e_abalone <- e
        best_reg_d_abalone <- d
        best_reg_model_abalone <- reg_model
      }
    }
  }
}
```

AVERAGE DISTANCE OF THE PREDICTED CLASS FROM THE TRUE CLASS

```
predicted_reg_abalone <- predict(best_reg_model_abalone, abalone_df[, - length(df)])
reg_abalone_distances <- abs(predicted_reg_abalone - abalone_df[, length(df)])
average_distance_reg_abalone <- mean(reg_abalone_distances)
```

HISTOGRAM OF FREQUENCY OF DISTANCES FROM THE TRUE CLASS

```
hist(reg_abalone_distances, main = "Histogram of Abalone Regression Distances",
     xlab = "Distance", ylab = "Frequency", xlim = c(0, 10))
axis(side = 1, at = seq(0, 10, 1), labels = seq(0, 10, 1))
```